

Consumer Intelligence

Operating Objectives for Pods towards GA · May 2026

JioBharatIQ

JioBharatIQ is India's daily AI Saathi — voice-first, multilingual, and built for everyday Bharat. The product spans health, devotion, entertainment, commerce, jobs, and more. The mission is to democratize intelligence for 1.4 billion Indians, making AI abundant, affordable, and accessible in every language, across every domain of daily life.

This document

This document sets the objectives for each pod in Consumer Intelligence. Every objective is numbered by priority, has a clear target, and carries a reason it matters. Objectives derive from the org-level GA Goals. Pod leads own every objective listed under their pod.a

What's covered in this document

AI Core: Voice & Multimodal · Agentic Orchestration + MLOps · Memory · Personalization · Evals · Quality · Guardrails

Experiences: Core Experience · Vertical Experiences · Partner Experience

Systems: Platform & Infrastructure · Data & Analytics

Insights & Growth: UXR · Consumer Lifecycle · GTM

AI CORE

Voice & Multimodal

Lead: Anmol Batra

Build and run the full non-text input/output pipeline. Voice covers audio capture, VAD, turn detection, barge-in, STT, language and code-mix detection, and TTS. Multimodal covers image and document understanding. Deliver clean structured intent to Agentic Orchestration.

#	Theme	Objective	Why it matters
1	Language Coverage	12+ Indian languages in production with full pipeline support — STT, language detection, code-mix handling, and TTS — each validated end to end before launch.	Partial support is worse than none. A language that passes STT but fails TTS produces confident wrong outputs.
2	Listening Quality	WER meets benchmark per language and model. VAD turn-taking and noise suppression accuracy validated against a golden dataset. Barge-in latency from TTS cancel to STT initiation within target.	STT errors compound downstream — a mis-transcribed word changes intent classification, skill routing, and the final response. Broken barge-in destroys conversational rhythm.
3	Voice Latency	End-to-end latency targets met across listening, VAD, TTFB, and synthesis — benchmarked per language and provider, tracked on a single dashboard.	Latency was the most cited pain point from Alpha. Above 2 seconds, voice users abandon mid-session.
4	Speaking Quality	TTS naturalness MOS above 4.0 per production language, validated by blind native-speaker evaluation.	Poor TTS in a user's language is the clearest signal the product was not built for them.
5	Scale and Cost	Reliability targets met at 100K+ concurrent load. Cost per STT minute and TTS character within defined targets.	At Jio scale, per-query costs dominate unit economics. The voice pipeline must scale without proportional cost increase.

Agentic Orchestration + MLOps

Lead: Bharat Bheesetti

Build and run the intelligence layer that turns user intent into completed tasks — routing, execution, cost discipline, and the MLOps platform that every pod depends on to deploy and operate models in production.

#	Theme	Objective	Why it matters
---	-------	-----------	----------------

1	Latency	Voice p50: simple ≤2.5s, single skill ≤7s, multi-skill ≤15s. Chat p50: simple ≤1.5s, single skill ≤4s, multi-skill ≤10s. First byte at ~30% of final latency on each class. One dashboard tracking all query classes.	Above 2 seconds on voice, users abandon mid-session. Chat users notice every 500ms beyond one second. Latency is where every other pod's work becomes visible to the user.
2	Routing Accuracy	97%+ of queries routed to the correct skill and tools at GA, measured weekly on a 1,100-query dataset. SLO holds without degradation as skill count crosses 100.	Wrong routing is invisible to the user — they receive a wrong answer with no signal about why. Accuracy must hold as the skill catalog grows.
3	Task Completion	95%+ of tasks initiated reach a resolved state. Hallucination ≤0.1%. Every irreversible action completes only after explicit user consent, with zero violations in monthly audit.	Routing correctly and completing the task are different problems. Irreversible actions at Jio scale require zero-tolerance consent enforcement.
4	Cost Efficiency	Weighted cost ≤Rs. 0.05 per query at projected GA QPS. 60%+ of queries served without a coordinator LLM call. Token cost p50 ≤3K.	The product cannot reach tier-2/3 economics without tight cost discipline. The pod owns the levers: cache hit rate, model tier mix, and tokens per query.
5	MLOps Platform	One model gateway, one registry, three in-house SLMs in production at GA: routing classifier, language and code-mix detector, simple-knowledge SLM. PR-to-prod median ≤30 minutes. 100% of orchestration config versioned, zero console edits.	Without a shared platform, every pod builds a shadow stack with no cost or eval discipline. In-house SLMs deliver Hinglish parity that no external model is trained for at the depth required.

Memory

Lead: Karan Prasad

Build and run the systems that capture, store, and retrieve what JioBharatIQ knows about a user. Every interaction should leave the product knowing the user better than before.

#	Theme	Objective	Why it matters
1	Write Quality	>80% extraction recall and >85% extraction precision on the golden eval dataset.	If the right memories are not written, nothing downstream works. Retrieval, personalization, and re-engagement all depend on a profile that reflects the actual user.
2	Read Quality	>80% Precision@K and >80% Recall@K on retrieval.	Retrieving the wrong memory is worse than retrieving nothing. An incorrect memory injected into a response degrades quality and erodes trust faster than a generic answer.
3	Latency	<500ms added to response latency at P95.	Memory must make responses better without making them slower. Slow or unavailable memory retrieval makes every interaction worse.
4	User Trust and Control	Every user can view, edit, and delete any stored memory. Deletions propagate across all surfaces within 24 hours. Zero deleted memories resurface.	Personalization only works if users trust it. Control is the condition under which users allow memory to work at all.
5	Continuous Optimization	Automated experimentation pipeline live. Every release candidate benchmarked against LongMemEval, BEAM, and JioEval before deployment.	Memory quality directly impacts personalization and task continuity. Continuous benchmarking is how quality compounds over time rather than drifting.

Personalization

Lead: Karan Prasad

Build and run the systems that make JioBharatIQ's responses and behavior adapt to each user — using memory, ecosystem signals, and real-time session context to shape what the product does for each person.

#	Theme	Objective	Why it matters
1	Day 0 Warm Start	>70% of new users have a populated personalization profile with name, language, and at least 3 ecosystem	Cold start is the highest drop-off risk. If pre-loaded user context does not fire on

		attributes loaded from JioID before their first in-app interaction.	session 1, the structural advantage of ecosystem signals is wasted before the user has typed a word.
2	Profile Completeness	>80% of users with 3+ sessions have a personalization profile with 10+ attributes populated — from JioID, session behavior, or memory — before their third session.	Thin profiles default to generic responses. The Jio data advantage is only real if it is populated and active early enough to shape the experience.
3	Signal Freshness	Personalization profiles reflect the latest user signals within 15 minutes of a session ending — validated daily.	A profile that lags by hours reflects who the user was, not who they are. Freshness is what makes personalization feel alive rather than stale.
4	Inference Latency	Personalization context injected into every eligible query within 50ms at P95.	Personalization sits on the critical path of every response. Any latency added here appears in every single interaction.
5	Cross-Surface Consistency	Any preference set on one surface propagates to all surfaces within 24 hours for >99% of users.	Surface-local contradictions — a language set in chat not respected in push notifications — directly undermine trust and are a DPDP compliance requirement.

Evals

Lead: Karan Prasad

Build and run the evaluation infrastructure that every pod depends on — CI/CD gates, eval suites, drift detection, and the tooling that makes quality measurement fast, self-serve, and adopted across the org.

#	Theme	Objective	Why it matters
1	CI/CD Coverage	Eval gates live on 100% of components, 1P verticals, and partner skills before any merge. Online-offline score delta <10pp at all times.	Any component change without CI/CD coverage can regress silently. 2P/3P skills have no internal deployment review — the eval gate is the only quality check before they reach users.
2	Continuous Monitoring	Online evals running on live traffic across all components, verticals, and partner skills. Drift detected within 24 hours, alert fires on any >10pp breach.	Regression prevention catches issues before go-live. Monitoring catches drift after. Without it, a model or data change can degrade live performance for days before anyone notices.
3	Self-Serve Platform	Automated golden dataset creation. Median time to first eval suite ≤2 days for any pod. Pods create and extend eval suites independently without Evals pod involvement.	If creating evals requires Evals pod involvement, coverage gaps accumulate faster than they can be closed. Speed of setup determines whether teams use the system or work around it.
4	Adoption	100% of deployments across all pods gated by an automated eval before reaching production. Every pod has at least one active eval suite per skill or component type they own. Coverage and gate compliance published weekly across all pods.	The gate ensures nothing ships unvalidated. Per-pod coverage ensures evals are embedded in how each team builds — not treated as a one-time setup that gets ignored after launch. Weekly publishing creates visibility and accountability at the pod level.
5	Golden Dataset Health	Golden dataset representativeness monitored continuously. If online scores diverge significantly from offline, the dataset is flagged and refreshed within one sprint.	A stale golden dataset means the offline quality bar no longer reflects what users actually experience. Eval coverage is only as good as the dataset it measures against.

Quality

Lead: Bhushan Patil

Own the quality bar across JioBharatIQ — defining what good looks like, validating that experiences meet the bar before they reach users, and maintaining quality standards across languages and verticals.

#	Theme	Objective	Why it matters
---	-------	-----------	----------------

1	Quality Bar Definition	Documented quality thresholds for every component, vertical, and partner skill type — owned by Quality pod, reviewed and updated quarterly.	Without a defined bar, quality decisions are subjective. The bar must be explicit so that go/no-go decisions are made on shared criteria, not individual judgment.
2	Pre-Launch Validation	100% of signature experiences validated against the quality bar before GA go-live. Zero experiences ship without a quality sign-off.	Alpha shipped with quality gaps that were knowable before launch. Pre-launch validation is the gate that prevents known regressions from reaching users.
3	Response Quality	80%+ response satisfaction across all signature experiences at GA. No single domain below 70%.	Response satisfaction is the direct user signal on quality. It is the measure that tells us whether the quality bar is calibrated correctly.
4	Cross-Language Parity	Quality scores within 10pp across all production languages. Any language falling below the floor triggers a dedicated quality sprint.	A quality gap in any language is a retention problem. Users in that language experience a worse product — and leave.
5	Quality Feedback Loop	Quality findings from live traffic reviewed weekly. Top issues surfaced to pod leads within 48 hours of detection.	Quality is not a launch gate — it is an ongoing practice. The feedback loop between live signal and pod action is what keeps quality from drifting post-GA.

Guardrails

Lead: Karan Prasad

Build and run the systems that protect users and the platform — ensuring safe, compliant, and policy-consistent outputs across every interaction and every language.

#	Theme	Objective	Why it matters
1	Coverage	100% of guardrails defined by legal live at Beta. Research agent continuously monitoring for new requirements across regulatory changes and emerging harm patterns.	Legal defines the floor, not the ceiling. The harm landscape evolves — a static list is out of date the moment it is written.
2	Effectiveness	FNR <5% per category. Hard block FNR (CSAM, UAPA, caste slurs): 0%. FPR <2% overall; Hard block FPR ≤0.5%. Cross-language FNR variance ≤10pp.	Under-blocking creates legal and safety risk. Over-blocking destroys user trust. Both are product failures.
3	Availability	Guardrails service uptime ≥99.9%. Latency budget ≤200ms. Hard block categories fail closed if detection layer is unavailable.	Guardrails sit on the critical path of every request. Downtime means unscreened traffic reaches the model — availability is a safety requirement, not a performance target.
4	Continuous Optimization	Production traffic continuously sampled and fed into the dataset engine. Classifiers retrained on production data on a defined cadence. Defined process for guardrail deployments across product and legal.	A guardrail trained only on synthetic data drifts from real user behavior. Closing the loop from production misses back into training keeps classifiers calibrated as language, topics, and evasion patterns evolve.

EXPERIENCES

Core Experience

Lead: Sridhar Iruvanti

Design and build the experience every user touches — home screen, conversation interface, onboarding, engagement loops, and how personalization surfaces visually. This is the layer everything else plugs into.

#	Theme	Objective	Why it matters
1	First Session Value	Above 40% of first-time users complete at least one task in their first session.	D1 return is directly correlated with first-session task completion. No completed task in session one means near-zero D7 return.
2	Onboarding Completion	Above 80% of new users complete onboarding in under 3 minutes across all device tiers and network conditions.	Every additional onboarding screen adds measurable drop-off, especially on low-end devices and slow connections.

3	Organic Daily Return	Above 25% of monthly active users return daily without a push trigger.	DAU/MAU driven by push notifications masks real engagement. Organic return is the measure of genuine habit.
4	Experience Discovery	Above 50% of users engage with at least 2 signature experiences in their first week without being explicitly directed.	Alpha showed near-zero organic discovery of vertical experiences without prompting. The home screen must do this work.
5	Bharat Accessibility	Above 90% of usability tests with first-time smartphone users in tier-2 and tier-3 cities pass without English dependency or literacy barriers across all production languages.	Any English dependency in navigation, prompts, or error states creates a hard barrier for the primary target user.

Vertical Experiences

Lead: Naroo Krishnan

Build and run the domain-deep experiences that make JioBharatIQ irreplaceable in the moments of Indian daily life that matter most. This is where the product earns its differentiation.

#	Theme	Objective	Why it matters
1	Breadth	At least 15 signature experiences live at GA across at least 5 verticals. Astrology + Devotion, Entertainment + News, and Jobs, Careers & Skills commit to 10 in total — the remaining 5 are committed by emerging vertical leads as part of GA planning.	15 signature experiences across 5 verticals is the GA commitment. It is the minimum to demonstrate category depth, not just general capability.
2	Depth	Every live signature experience at GA is structurally unreplicable — requiring Jio data, licensed content, or ecosystem access no standalone AI can bring.	An experience that only requires a better prompt is not a moat. Structural unreplicability is what makes each experience defensible.
3	Retention	Every signature experience shows D7 organic return above 40%.	D7 organic return per experience is how the data flywheel starts. Daily users generate the signals that make personalization compound.
4	Velocity	Time from validated use case to live experience below 7 days across all active verticals — applied across the full L0 to L3 quality ladder. Vertical Lift, an active project within the team, is building the agentic pipeline that makes this possible at scale — agents carry the repeatable work across landscape, build, and iteration stages, with human judgment applied at the right points.	Build velocity determines how much of the roadmap ships by GA. Without an agentic pipeline, running many experiences in parallel across multiple verticals simultaneously is a throughput problem the team cannot solve manually. Vertical Lift is how breadth and quality happen at the same time.

Astrology + Devotion

Lead: Shivali Wason

Build the daily spiritual and astrological companion for Bharat — the screen where Indians start their day, explore their Kundali, attend live darshan, and stay connected to their faith across every tradition.

#	Theme	Objective	Why it matters
1	Adoption	≥30% of JioBharatIQ MAU engage with Astrology weekly. ≥40% with Devotion weekly. ≥60% of new users try each vertical within their first 7 days.	120M+ Indians engage with astrology and hundreds of millions with daily devotion through fragmented, low-quality sources. Adoption measures whether JioBharatIQ's version is compelling enough to pull them in.
2	Activation	≥80% of first-session astrology users leave knowing at least 3 things about their Kundali they did not know before. ≥70% of first-time devotion users complete a full devotional act without needing another app.	The first session is the only audition. If users leave with generic affirmations or still need a separate app for darshan, the vertical has not earned the next session.
3	Daily Habit	≥25% of Devotion MAU engage daily. Engaged astrology users return at least 4 days per week, organically. Session depth and return frequency trend upward from session 3 to session 10.	Devotion is structurally daily. If the daily return number is low, we built a notification, not a practice. Upward trend over sessions confirms the experience is compounding, not plateauing.

4	Quality and Trust	Above 85% of astrology interpretations rated accurate and personally relevant. NPS ≥60 after 3rd session for both verticals.	One inaccurate reading destroys the trust built over three sessions. NPS ≥60 means users recommend to family — in devotion and astrology, personal endorsement is the only acquisition channel that scales.
5	Iteration Velocity	Top 3 unmet user intents per vertical published weekly. New experiences and improvements shipped within 7 days of a validated signal.	The launch experiences are the best hypothesis, not the final answer. Users across faiths and astrological needs will tell us what is missing. 7-day cycle time is only possible because we are AI-first.

Entertainment + News

Lead: Shantanu Maheshwari

Build the AI companion for India's highest-attention daily moments — live cricket, music, morning news, and fact-checking in a country where misinformation travels faster than corrections.

#	Theme	Objective	Why it matters
1	Adoption	≥5% of JioHotstar match viewers actively use Cricket Companion during live matches. ≥5% of JioSaavn users stream music on JioBharatIQ weekly. ≥5% of JioNews users use JioBharatIQ for daily news.	Adoption from within the Jio ecosystem validates the structural advantage. These numbers confirm the ecosystem integration is working, not just that the product exists.
2	Activation	≥35% of new JioBharatIQ users try an Entertainment or News experience within their first 7 days.	Entertainment is the widest-reach vertical. High activation here seeds multi-vertical usage and drives DAU/MAU across the full product.
3	Engagement	≥25% of JioBharatIQ MAU engage with an Entertainment or News experience at least once a week.	Weekly engagement is the measure of whether these experiences are part of the user's regular rhythm — not just a one-time trial.
4	Quality	CSAT ≥4.0 for Cricket Companion, AI Jockey, and AI Anchor. Fact Checker resolves ≥90% of submitted WhatsApp forwards within 8 seconds with a sourced verdict.	A fast, sourced fact-check verdict is what no incumbent news or social platform delivers. Speed and sourcing together are what make Fact Checker structurally unreplicable.
5	Shareability	≥10% of Cricket sessions end in a WhatsApp share. ≥15% of Fact Checker verdicts shared back into WhatsApp.	Sharing is acquisition. In cricket and fact-check, every share draws a social group in. Share rate is the metric that tells us whether the content is earning organic distribution.

Jobs, Careers & Skills

Lead: Samyak Shah

Build the personal career and learning coach for India's 200M+ job seekers and exam aspirants — with real coaching and real feedback, not content without consequence.

#	Theme	Objective	Why it matters
1	Adoption	≥30% of JioBharatIQ MAU engage with Jobs, Careers & Skills weekly. ≥50% of new users try Jobs, Careers & Skills within their first 7 days. ≥40% of active Jobs, Careers & Skills users name JioBharatIQ as their primary destination for career prep, unprompted.	Jobs, Careers & Skills is not structurally daily like devotion — adoption is harder won. These numbers say Jobs, Careers & Skills is part of the weekly rhythm and that discovery and onboarding are routing users into it effectively.
2	Activation	≥50% of new Jobs, Careers & Skills users complete a first coached session — diagnostic answered, coaching delivered, feedback returned to the user.	A user who has not completed a first session has nothing to return to. The first session is the only place to prove that JioBharatIQ does what no free product does: deliver personal, specific feedback.
3	Coaching Quality	≥75% of started sessions end with a feedback signal — diagnostic, correction, or readiness score — delivered to the user.	If sessions do not end with feedback, the coaching premise has broken. Content delivered without feedback is a library, not a coach.

4	Habit and Return	≥40% of users who complete a first session return within 7 days. Engaged users sustain 3 to 4 active days per week, with at least 50% of repeat sessions organic.	Coaching has a different rhythm to devotion. 3 to 4 days per week is the honest bar that says Jobs, Careers & Skills has earned a regular slot in the user's week. Organic return confirms the habit is real.
5	Outcomes	≥20% of users with a logged real-world outcome event report a positive outcome attributed to Jobs, Careers & Skills within 90 days. NPS ≥55 after 3rd session.	Jobs, Careers & Skills is the only JioBharatIQ vertical where the user has a measurable real-world result. Outcome attribution is the proof that the coaching loop is closing. NPS ≥55 in a vertical driven by aspiration means users recommend it to people they care about.

Partner Experience

Lead: Sushant Joshi

Build and run the Partner Portal — the infrastructure that lets partners onboard, test, publish, and manage their skills on JioBharatIQ. Partner relationships and integration work are owned by the vertical leads within their domains.

#	Theme	Objective	Why it matters
1	Onboarding Speed	A new partner reaches a working test skill in sandbox within 30 minutes of sign-up.	Fast time-to-first-working-skill converts partner interest into commitment. Friction at setup is the leading cause of partner drop-off before any skill is submitted.
2	Skill Quality Gates	80% of governance checks automated by GA — no skill that fails the quality bar reaches a user because a human reviewer was the last line of defense.	A platform is only as trustworthy as its worst live skill. Automated quality gates are what allow the catalog to scale without scaling the review team linearly.
3	Integration Range	All three integration types live before GA — MCP, API, and Agent — so no partner is turned away by a technical ceiling.	The integration type a partner needs is determined by their stack, not ours. If we do not support it, they do not join.
4	Partner Visibility	Every skill's behavior visible every 5 minutes. Errors and policy violations published within 1 minute.	Partners cannot improve what they cannot see. Real-time visibility is what turns a live skill into a managed product rather than a black box.
5	Self-Serve Operations	Partners update, version, and improve their live skills in 9 out of 10 cases without platform team involvement.	At scale, the platform team cannot be the critical path for every partner action. Self-serve is what allows the partner ecosystem to grow beyond a managed cohort.

SYSTEMS

Platform & Infrastructure

Lead: TBD, Manan Verma, Ayush Sharma

Build and run the shared infrastructure every pod depends on — services, reliability, delivery pipeline, and the systems that make every pod's work deployable, observable, and scalable.

#	Theme	Objective	Why it matters
1	Availability and Reliability	JioBharatIQ available above 99.5% at GA. Incidents detected within 3 minutes, acknowledged within 5 minutes. Full observability across 100% of production services with log-trace correlation. MTTR <2 minutes, MTTR <30 minutes.	99.5% availability at 1M+ users means under 4 hours of downtime per month. You cannot operate what you cannot see — unified observability is the difference between a 5-minute triage and a 2-hour outage.
2	Scale	Platform load-tested and stable at 100K+ concurrent users with 30% headroom before GA. Production-grade GKE with HPA, KEDA, and Cluster Autoscaler. Zero pods Pending >3 minutes. Zero downtime on upgrades.	A load test failure discovered post-GA is a public incident. Without autoscaling from day one, every traffic spike is an incident waiting to happen.
3	Latency SLOs	Every service holds its latency SLO at 100K concurrent load — voice below 1,500ms p95, agentic below 700ms p95.	Platform latency SLOs are shared dependencies. A breach means Voice, Agentic, and Memory pods all miss their targets simultaneously.

4	Delivery Velocity	Every service ships to production via a standard GitOps pipeline. Image-to-prod under 10 minutes. Any pod ships from PR to production in under 2 hours, gated by automated evals. Manual deployments fully retired.	Standard GitOps turns every release into a reviewable, auditable commit. Under 2 hours means pods can ship, observe, and iterate same-day.
5	Service Autonomy	Any pod stands up a new service with CI/CD, monitoring, and eval gating in under 2 days without a platform ticket. All infrastructure provisioned via Terraform and deployed via ADO.	Every new service that needs a platform ticket adds days of wait time for the requesting pod. IaC keeps governance, security, and cost discipline consistent across every service.

Data & Analytics

Lead: Kalyan Datta

Build and run the data infrastructure, analytics systems, and experimentation tooling that let every pod measure what matters, move fast on evidence, and trust the numbers they are looking at.

#	Theme	Objective	Why it matters
1	Metric Coverage	Every objective in this document has a live, accurate dashboard before GA launch.	An objective without a live dashboard is unmanageable at GA scale. Pods need to see their numbers daily, not on request.
2	Data Freshness	All data sources joined on user and session within 5 minutes — validated daily.	A 24-hour data lag means decisions are made on data that is already one iteration cycle old.
3	Self-Serve Analytics	Any team member answers a product question or builds a dashboard in under 2 minutes without SQL or data engineer involvement.	Questions that need a data engineer get deprioritized or skipped. Self-serve access is what makes evidence-driven decisions the default.
4	Instrumentation Velocity	Adding a new data point to any consumer takes under 1 day.	New features with no instrumentation ship into a black box. A 1-day turnaround keeps measurement current with the build.
5	Experimentation Infrastructure	GrowthBook live at GA — any pod can run an A/B test end to end in under 3 days without data engineering involvement.	At 1M users, the cost of a wrong product decision is large. A/B testing must be accessible to every pod without specialist support.

INSIGHTS & GROWTH

UXR

Lead: Muz Mukhtar

Generate the user understanding that shapes what JioBharatIQ builds and how well it works — from foundational research on who our users are, to validation before ship, to measurement after launch.

#	Theme	Objective	Why it matters
1	Foundational Coverage	Validated cohort definitions, day-in-life archetypes, pain points, and opportunities for all primary user cohorts published and adopted by product and design teams by Beta exit.	Without validated personas, product decisions default to the team's own profile — urban, English-literate, tech-native. That is not the target user.
2	Pre-Ship Validation	Above 90% of signature experiences and core experience changes have a UXR validation study completed before GA launch, with findings shared at least 2 weeks before the ship date.	Alpha surfaced latency issues and near-zero organic vertical discovery — both were knowable before launch. Pre-ship validation catches these before they reach 1M users.
3	Post-Launch Sentiment	Structured sentiment pulse live within 2 weeks of GA — CSAT and top-3 pain point signal across all active user cohorts, delivered to product and design leads within 2 weeks of collection.	GA is when real users hit the product at scale for the first time. Without a structured post-launch loop, negative experiences compound silently.
4	Research Infrastructure	Panel, survey tooling, and measurement framework live and producing data on at least 5 active studies before GA.	Qualitative research does not scale to 1M users. Quant infrastructure is what makes UXR relevant at GA.
5	Research Velocity	Time from research question to actionable insight below 5 days for iterative studies and below 21 days for	Insight delivered after a build decision is documentation, not research. Turnaround

		foundational studies, measured across all studies in the 60 days before GA.	time determines whether findings change anything.
--	--	---	---

Consumer Lifecycle

Lead: Ayush Prashant

Own the signals, interventions, and experimentation that drive activation, retention, and re-engagement. CLM does not own the experience — product and vertical PMs own what the user sees. CLM owns the lifecycle layer that operates on top of it: when to reach users, how, and with what signal.

#	Theme	Objective	Why it matters
1	Same-Session Activation	Above 40% of acquired users complete at least one task within their first session, across all acquisition channels.	Same-session activation is the leading indicator of downstream repeat. An installed but inactive user inflates MAU while suppressing DAU/MAU.
2	Early Retention	DAU/MAU ratio of 20%+ and D30 retention rate of 30%+ at GA.	DAU/MAU is the measure of whether users build a daily habit or inflate the install count. Below 20%, the MAU target is a number that decays within weeks of launch.
3	Engagement Depth	Multi-domain engagement rate increasing month-over-month — users active across 2+ use case domains within 30 days of first session.	Multi-domain usage compounds retention. A user anchored to one vertical is more likely to churn than one who has found value across the product.
4	Decay Capture	At-risk users identified and re-engagement campaigns triggered before they become dormant. Recovery rate above 15% of lapsed users within 30 days.	At scale, saving 5% of at-risk users moves thousands of DAU. Without proactive signals, churn is only visible after it has already happened.
5	Experimentation Engine	Experimentation infrastructure live and end-to-end tested in FnF/Beta — hypothesis to assignment to measurement to readout. Any pod can run a lifecycle experiment without CLM involvement.	Learning velocity determines how fast lifecycle interventions improve. Without a working loop, every lifecycle decision is a guess rather than a tested hypothesis.

GTM

Lead: Paarmi Modi

Drive user acquisition and launch execution to reach 10M+ MAU within 90 days of GA.

#	Theme	Objective	Why it matters
1	Cohort Identification	Primary acquisition cohorts for GA identified and ranked by conversion opportunity, ease of reach, long-term strategic value, and acquisition cost — finalized before channel planning begins.	The cohort framework determines channel mix, spend sequencing, and what success looks like. Going after the wrong cohort first wastes spend and generates low-quality retention signal.
2	Brand and Positioning	Brand identity, consumer positioning narrative, and GA comms approach locked with leadership sign-off before any creative production timelines start.	Every channel, creative, and partner brief must reflect a consistent narrative. Locking early prevents downstream rework across every workstream.
3	Acquisition Execution	10M+ MAU within 90 days of GA. Channel mix sequenced across Jio ecosystem, BTL, social, and creator programs — with a clear view of what is live at GA and what follows.	Jio distribution is the structural acquisition advantage. A clear channel plan is what converts that advantage into activated users.
4	Launch Program	Cross-functional launch program with a signed-off plan, go/no-go checklist grounded in Beta signal, and all legal, policy, and crisis-comms approvals cleared before GA.	Without agreed gate criteria and governance sign-off, launch decisions become subjective. Missing approvals can halt the launch.