

Consumer Intelligence

GA Plan · May 2026

What each pod is building, and in what order. Beta comes first, then GA. The plan is derived from pod leads' own planning and logically ordered where sequencing was not explicitly specified.

AI CORE

Voice & Multimodal

Beta The team establishes listening quality baselines — golden datasets across languages, code-mixing, and noise environments — and gets VAD accuracy, noise suppression, and MOS evaluation running. Language coverage expands from 7 to 12, with full STT, code-mix handling, and TTS for each. The pipeline gets hardened with response language and script matching across the full stack, rich UX card support, and image and document guardrails. An in-house language detection model improves ASR routing accuracy.

GA TTS voice selection is finalized per language, prosody tuned per content type, and multimodal understanding validated on core Indian use cases — bills, medicine labels, invoices, IDs. The pipeline is load-tested at 100K+ concurrent users and intelligent model routing with per-language fallbacks ships to production.

Agentic Orchestration + MLOps

Beta Routing accuracy is benchmarked weekly on a 1,100-query golden dataset targeting 97%+. Latency SLOs are enforced across all query classes on a single dashboard. A semantic cache layer targets 60%+ of queries served without a coordinator LLM call. The human-in-the-loop architecture for irreversible actions is built and tested. CI/CD eval gates are live on all orchestration components.

GA Three in-house SLMs are in production — routing classifier, language and code-mix detector, and simple-knowledge SLM. The platform holds routing accuracy and latency SLOs at 100K+ concurrent load and 100+ skills. Cost per query is at or below Rs. 0.05. PR-to-prod median is under 30 minutes, with 100% of config versioned and zero console edits.

Memory

Beta Two structured experiments each for extraction precision/recall and retrieval Precision@K/Recall@K move the pod toward its 80%/85% targets. Component-level eval pipelines run offline against the golden dataset. Memory Arena migrates to Kubernetes for continuous benchmarking on LongMemEval, BEAM, and JioEval. User-facing memory management — view, edit, delete — ships with propagation across surfaces within 24 hours. Extracted memories are pushed to OCP for personalization consumption.

GA Quality metrics and latency targets hold at P95. The experimentation pipeline runs automated for ongoing optimization.

Personalization

Beta Phase 1 of the lookalike model computes cohort affinity vectors daily from high-engagement users and incorporates UCP signals for cold-start personalization. Phase 2 extends coverage to sparse users with 1-20 sessions. A CleverTap data pusher sends personalized signals from OCP for push notification and WhatsApp dispatch. APM metrics and functional domain events are instrumented throughout.

GA The personalization pipeline serves dynamic rankings for home and vertical pages. Cross-surface consistency is enforced and inference latency holds within 50ms at P95. The experience is measurably better at session 10 than session 1.

Evals

Beta CI/CD gating is wired into the DevOps pipeline and returns pass/fail on every PR merge. Eval suites are built for 8 safety categories currently uncovered. An automated golden dataset creation pipeline feeds grader creation and keeps datasets fresh. Phase 1 of the evals dashboard is live covering online scores, offline regression history, and grader calibration. A guardrails research agent monitors regulatory changes and emerging harm patterns continuously.

GA Partner skill go-live gating is enforced. Any pod deploys a new eval suite in under 2 days without Evals pod involvement. A competitive benchmarking tool runs JioBharatIQ and competitors against the same golden dataset. Per-pod adoption — 100% of deployments gated — is tracked and published weekly.

Quality

Beta Quality thresholds are documented for every component, vertical, and partner skill type. Every signature experience is validated against the bar before go-live. Cross-language quality scores are tracked, and any language falling below 70% satisfaction triggers a dedicated sprint.

GA Response satisfaction reaches 80%+ across all signature experiences with no single domain below 70%. A weekly quality feedback loop surfaces top issues from live traffic to pod leads within 48 hours of detection.

Guardrails

Beta Tier 1 guardrails cover child safety (PhotoDNA, Thorn), PII (policy and consent engine), and hate speech (MuRIL classifier). Prompt-level guardrails are live for five categories — financial advice, medical advice, IP, AI disclosure, and dark patterns. FNR instrumentation compares eval grader scores against guardrail actions to surface false negatives. Hard block categories fail closed when the detection layer is unavailable.

GA Five additional MuRIL classifiers cover terrorism, adult content, national symbols, electoral manipulation, and financial/medical categories. An adversarial dataset spans all 29 guardrail categories. A production traffic sampling pipeline feeds confirmed misses back into classifier retraining on a defined cadence. A guardrails observability dashboard tracks FNR, FPR, and latency per guardrail.

EXPERIENCES

Core Experience

Beta Context-aware onboarding eliminates unnecessary screens by inferring device language, time, and location at launch. JioID detected users skip straight to home. The allowlisting and volume gate mechanism manages cohort-based rollout with a kill switch. GrowthBook feature flags enable any feature to be toggled without an app update. Home page dynamic ranking consumes the Personalisation API for both first-time and returning users. The i18n framework enables Tamil, Telugu, Gujarati, and Kannada with full UI in Beta. WebSocket transport and the analytics event framework are foundational infrastructure the rest of the product depends on.

GA JioID native SDK integration replaces the existing integration. Performance across cold start, render time, memory footprint, and bundle size meets targets on low-end devices. A full zero-English dependency audit clears all navigation, prompts, and error states across all production languages. Every core flow is completable without a keyboard. Crash, ANR, and OOM monitoring goes live. Automated notification orchestration triggers all campaigns without manual configuration.

Vertical Experiences

Beta Vertical Lift POC replicates 1-2 Astrology experiences to validate the agentic build approach. Reusable primitives — intent sets, evaluation harnesses, and build scaffolding — are identified and documented across verticals. The L0 to L3 quality ladder is defined per vertical and every active experience is tracked against it weekly.

GA 15 signature experiences go live across 5+ verticals. Every live experience targets D7 organic return above 40%. The Vertical Lift pipeline drives parallel builds across verticals with a 7-day cycle time from validated use case to live, tracked and published per vertical.

Astrology + Devotion

What ships *Grand Kundali Reveal · Talk to Your Planets · Kundali Matching · Shubh Muhurat Finder · Live Darshan · Devotional Feed · Sacred Calendar*

Beta The foundational elements are already live in FnF — Kundali creation, Shubh Muhurat, and Kundali Matching. Beta is about deepening them: connecting the four Astrology experiences into a coherent journey and building the Devotion content engine.

GA Astrology ships Grand Kundali Reveal (the cinematic voice-first walkthrough), conversational planetary exploration with session memory, Guna Milan compatibility analysis, and personalized auspicious timing. Devotion ships live darshan with worship booking and Prasad delivery, a daily devotional content feed shareable to WhatsApp, and a sacred calendar covering all faith traditions with vrat katha, fasting guides, and cross-faith observances.

Entertainment + News

What ships *Cricket Match Companion · AI Jockey · AI Anchor Daily Brief · Multimodal Fact Checker*

Beta Beta builds the data foundations and core experiences. For Cricket: JioHotstar live match data feed is integrated, the prediction engine ships with friend-group leaderboard, AI voice commentary goes live in Hindi and English, and a voice query engine answers natural language match questions in real time. For Music: JioSaavn integration enables mood-to-playlist in one sentence with an AI RJ base persona. For News: the signal ingestion pipeline, morning brief generator, and profile-aware item ranking are live. For Fact-Check: text claim intake, the four-tier verdict engine, source display, image OCR, and voice note transcription all ship in Beta.

GA Cricket commentary expands to Tamil, Telugu, and Marathi. Match summaries become shareable as WA cards. Music taste memory extends to 30 days with the AI RJ referencing listening history. News briefs launch in Marathi, Tamil, Telugu, and Bengali. The WhatsApp bot for fact-check lets users forward directly to JioBharatiQ without opening the app. Fact-check verdict cards are WA-optimised and share-back rate is tracked against the 15% target.

Jobs, Careers & Skills

What ships *Microlearning · English Learning · Interview Prep · Government Exam Discovery & Prep*

Beta The diagnostic-first session opener — the shared primitive across all four use cases — is defined and the build approach is locked. Vertical Lift is used to accelerate parallel development across microlearning, English learning, interview prep, and government exam prep.

GA Microlearning delivers problem-led sessions with same-day usable outputs. English learning provides real-time pronunciation and tone feedback with a conversational AI avatar that remembers across sessions. Interview prep delivers company-specific mock studios with pressure drills and a morning-of kit. Government exam discovery maps aspirations to exam paths before naming any exam, with adaptive prep and commute-orchestrated daily sessions. Outcome instrumentation runs from day 1 — every user who logs an interview, exam, or English event gets a 2-tap follow-up.

Partner Experience

Beta A sandbox environment and playground let partners test skills across text and voice before submitting. Prompt governance, PII validation, T&C acceptance, and admin review tooling form the quality gate before any skill goes live. Skill versioning ensures every change goes through the governance cycle. A partner portfolio dashboard gives partners a single view of all their skills and states. Developer documentation covers all integration types.

GA An execution details panel in the playground surfaces Langfuse traces. A risk-tier publish workflow routes submissions by data sensitivity and governance score, fast-tracking low-risk skills. The portal home page serves multiple partner states. Governance moves from advisory to hard gate enforcement. Skill performance alerts notify partners automatically when a skill degrades.

SYSTEMS

Platform & Infrastructure

Beta Not applicable — the Platform & Infrastructure team works to a GA milestone. The work is sequenced as a single pre-GA program.

GA Services move from manual kubectl to a standard GitOps pipeline via ADO, ArgoCD, and Helm. Image-to-prod targets under 10 minutes. Full observability lands on New Relic — APM, infra, logs, and distributed tracing across 100% of production services, with MTTD under 2 minutes and runbooks attached to every P0 alert. GKE runs production-grade with HPA, KEDA, and Cluster Autoscaler — zero pods Pending over 3 minutes, zero downtime on upgrades. Databricks on GCP with Unity Catalog provisions the data platform. The GCP External HTTPS Load Balancer with CDN is load-tested to 100K+ concurrent users with p99 under 500ms before launch.

Data & Analytics

Beta Every pod objective has a live, accurate dashboard before Beta launch. Data freshness is validated at 5-minute join latency on user and session. GrowthBook experimentation infrastructure is live and any pod can run an A/B test end to end in under 3 days. Self-serve analytics lets any team member answer a product question in under 2 minutes without SQL.

GA All production data pipelines get automated quality checks with anomalies surfaced within 30 minutes. Adding a new data point to any consumer targets under 1 day.

INSIGHTS & GROWTH

UXR

Beta Cohort hypotheses and segmentation are defined from secondary research and approved by leadership. Qualitative diary studies and IDs produce validated persona definitions and day-in-life archetypes per cohort. These are published to the knowledge base and formally adopted by product and design teams by Beta exit. The pre-ship validation engine playbook is authored with Consumer Lifecycle, setting the standard brief, recruitment criteria, and reporting template for all pre-ship studies.

GA Foundational research on Astrology, Devotion, News, Entertainment, and Jobs/Careers/Skills produces 5 key insights per vertical, built in partnership with vertical leads. 90%+ of signature experiences have a UXR validation study completed at least 2 weeks before their ship date. A structured sentiment pulse goes live within 2 weeks of GA, delivering CSAT and top-3 pain point signal across all active user cohorts to product and design leads.

Consumer Lifecycle

Beta Activation event taxonomy is defined per vertical with a dashboard tracking same-session activation rate. The homepage maps to 6 lifecycle states — New, Activated, Engaged, Retained, At-Risk, and Returning — each with defined content behaviour. A D1/D3/D7 nudge sequence delivers value callbacks and adjacent use-case discovery across push notification and WhatsApp. The lifecycle state machine is implemented in a 7-state model synced to CleverTap daily. At-risk users are identified daily and pushed to a CleverTap segment for automated winback triggers. GrowthBook experiment workspace is live with a 10% global holdout.

GA Daily loops launch for Devotion, Astrology, and News. At-risk users enter a 3-touch winback sequence. Channel fallback logic routes across WhatsApp, push notification, and in-app. A consolidated intervention calendar prevents user over-messaging. Experiment decision velocity — ship, kill, or iterate — is tracked from Beta and scales at GA.

GTM

Beta Primary acquisition cohorts are identified and ranked by conversion opportunity, ease of reach, strategic value, and acquisition cost. Brand identity and consumer positioning narrative are signed off by leadership before any creative production begins. The MarTech stack — AppsFlyer attribution, agency partners across creative, media, and PR — is contracted and live.

GA The acquisition channel mix across Jio ecosystem, BTL, social, and creator programs is sequenced and ready at launch. A cross-functional launch program with a signed-off go/no-go checklist and all legal and crisis-comms approvals cleared runs before launch. The target: 10M+ MAU within 90 days of GA, with D30 retention at 30%+ and DAU/MAU at 20%+.